

Black–Litterman with Ledoit–Wolf Shrinkage: Fixed- Ω Leakage and the Critical Correlation Threshold*

Lucas Posern

Faculty of Mathematics, TU Dresden

posernlucas@gmail.com

July 14, 2026

Abstract

The Black–Litterman model is often combined with shrinkage estimators of the covariance matrix, most prominently the Ledoit–Wolf estimator. We show that the effect of the estimator switch on the posterior depends on how the view-uncertainty matrix Ω is handled. Under Idzorek’s confidence calibration, Ω co-moves with Σ and, for a single view, the switch is exactly inert in the view channel; with several views the deviation is driven by the cross-view covariances and, in our three-view application, stays an order of magnitude below the effects studied here. If instead Ω is calibrated once on the sample covariance and then frozen, a workflow that can easily arise when the covariance model is upgraded after the views have been signed off, the effective view confidence c_{eff} shifts by a closed-form amount that is independent of the prior scale τ and grows with the shrinkage intensity. The shift is positive precisely when the view variance lies below twice the average asset variance, a condition met by relative views on highly correlated pairs *and* on low-variance (bond) pairs; for a general view vector the threshold is $\mu_s \|p\|^2$, and on our data an absolute view on the lowest-variance asset shifts by +12.9 percentage points, more than any pair. We derive an exact expression for the induced changes in posterior means and weights under a single relative view, provide a closed-form critical correlation threshold, and quantify the effect on an ETF implementation of Idzorek’s eight-asset universe: at a shrinkage intensity of 0.041, five of 28 pairs shift their effective confidence by more than five percentage points, the largest by +10.7. Recomputing Ω alongside Σ removes the leakage exactly.

Keywords: Black–Litterman; Ledoit–Wolf shrinkage; view uncertainty; covariance estimation; portfolio construction.

JEL: G11; C13; C58.

1 Introduction

The Black–Litterman model (Black and Litterman, 1992) is among the most widely implemented models of quantitative portfolio construction. It anchors expected returns at the CAPM market equilibrium (Sharpe, 1964) and blends in subjective views with an explicit uncertainty matrix Ω , thereby taming the input sensitivity of mean–variance optimisation (Markowitz, 1952) documented by Michaud (1989), Best and Grauer (1991) and Chopra and Ziemba (1993); Bayesian shrinkage of the mean (Jorion, 1986; Frost and Savarino, 1986) and decision-theoretic treatments of parameter uncertainty (Kan and Zhou, 2007) attack the same problem from the estimation side. Its second input, the covariance matrix Σ , receives comparatively little attention in the

*This paper grew out of the author’s bachelor thesis at TU Dresden. I thank Rodrigo S. Targino (FGV EMap, Rio de Janeiro) for detailed feedback on the proof strategy. All remaining errors are mine.

Black–Litterman literature (empirical exceptions include [Bessler et al., 2017](#)), although the sample estimator is unreliable whenever the number of assets is not small relative to the sample size. A standard remedy is the shrinkage estimator of [Ledoit and Wolf \(2004a\)](#), which pulls the sample covariance toward a scaled identity and improves the conditioning of Σ^{-1} substantially: on our data the condition number falls by a factor of three, and by far more as n approaches T ([Ledoit and Wolf, 2004b](#)).

This paper studies what happens inside the Black–Litterman model when the covariance estimator is switched from the sample estimator to Ledoit–Wolf. Within the view channel, the answer turns on a bookkeeping question that, to our knowledge, has not been isolated in the literature: the treatment of Ω . Two workflows arise naturally. In the first, Ω is derived from Σ through Idzorek’s confidence calibration ([Idzorek, 2007](#)) and therefore co-moves with the estimator. In the second, ω_k is computed once on the sample covariance, stored as a set of scalars, and left untouched when the covariance model is later upgraded. We call the resulting shift of the effective view confidence away from the stated confidence, at frozen Ω , *leakage*.

Our contribution is threefold. First, we show that under Idzorek’s calibration the estimator switch is *exactly* inert in the view channel for a single view: the effective view absorption equals the specified confidence c for any positive-definite estimator, and the remaining effect on the posterior runs through the equilibrium prior, which is exactly linear in the shrinkage intensity α , and the view projection, which is linear in α to first order. With several simultaneous views the deviation is driven by the cross-view covariances; Remark 1 quantifies it. Second, for the fixed- Ω workflow we derive a closed-form expression for the shift Δc_{eff} of the effective view confidence. The expression is exact, independent of τ (the same τ enters calibration and posterior, and cancels), linear in the shrinkage intensity α for small α , and its sign is determined by a simple criterion: the view variance versus twice the average asset variance. In the symmetric case of equal asset volatilities this criterion translates into a closed-form critical correlation threshold $\rho^*(\alpha, \varepsilon)$. An exact identity for the Black–Litterman posterior under a single relative view then propagates Δc_{eff} analytically into posterior means and portfolio weights. Third, we quantify the effect on an ETF implementation of the eight asset classes in [Idzorek \(2007\)](#): with the analytic shrinkage intensity $\alpha = 0.041$, five of the 28 asset pairs shift their effective confidence by more than five percentage points; among them is a bond pair with moderate correlation, which shows that high correlation is not necessary for a material shift. Separating the view channel from the precision channel through Σ^{-1} shows that a frozen Ω restores roughly 40 % of the gross-exposure reduction that the estimator switch delivers.

Covariance estimator, confidence and view uncertainty form one unit; upgrading one piece without re-deriving the others moves the effective model away from the analyst’s stated inputs. Related observations (that constraints or estimators interact with estimation error in ways invisible at the input stage) appear in [Jagannathan and Ma \(2003\)](#) for long-only constraints and in [DeMiguel et al. \(2009\)](#) for the difficulty of beating naive diversification out of sample. Within the Black–Litterman literature, [He and Litterman \(1999\)](#), [Satchell and Scowcroft \(2000\)](#), [Herold \(2003\)](#) and [Walters \(2014\)](#) document the model’s sensitivity to the (τ, Ω) specification (see [Meucci, 2010](#) for a survey, [Kolm and Ritter, 2017](#) for the Bayesian structure, [Michaud et al., 2013](#) for a critical view, and [Bevan and Winkelmann, 1998](#) for a production workflow); our result gives one mechanism of this sensitivity a closed form. On the shrinkage side, [Ledoit and Wolf \(2003\)](#), the nonlinear extensions of [Ledoit and Wolf \(2017\)](#) and [Ledoit and Wolf \(2020\)](#), and the portfolio-oriented intensity calibration of [DeMiguel et al. \(2013\)](#) would enter our formulas only through a different expression for the view-variance shift; the leakage mechanism itself is target-independent.

Section 2 fixes notation, model and data. Sections 3 and 4 review the conditioning problem and the Ledoit–Wolf estimator. Section 5 settles the Idzorek case. Section 6 contains the main results for fixed Ω . Section 7 isolates the precision channel through Σ^{-1} and quantifies the weight-level cost of the frozen- Ω workflow; Section 8 concludes.

2 The Black–Litterman Model and the Data

For n risky assets with covariance $\Sigma \in \mathbb{R}^{n \times n}$, $\Sigma \succ 0$, the Black–Litterman posterior mean is

$$\mu_{\text{BL}} = [(\tau\Sigma)^{-1} + P^\top \Omega^{-1} P]^{-1} [(\tau\Sigma)^{-1} \pi + P^\top \Omega^{-1} Q], \quad (1)$$

with implied equilibrium return $\pi = \delta \Sigma w_{\text{eq}}$, view-pick matrix $P \in \mathbb{R}^{K \times n}$, view returns $Q \in \mathbb{R}^K$, view-uncertainty matrix $\Omega \succ 0$, prior scale $\tau > 0$, and risk aversion $\delta > 0$. The unconstrained portfolio is $w_{\text{BL}} = \frac{1}{\delta} \Sigma^{-1} \mu_{\text{BL}}$.¹

Data. We use an ETF implementation of the eight asset classes in Idzorek (2007): AGG (US bonds), BWX (international bonds), IWF/IWD (US large-cap growth/value), IWO/IWN (US small-cap growth/value), EFA (international developed equity) and EEM (emerging markets equity). We take monthly total returns from Yahoo Finance (adjusted closes) and convert them to excess returns with the 13-week T-bill rate; the series are complete over the whole analysis period. All estimates use the window January 2008 to December 2018 ($T = 132$, $n = 8$). The market weights w_{eq} are the market-capitalisation weights published by Idzorek for these asset classes; they predate the sample (2004 vintage), which makes them a dated but look-ahead-free prior anchor. All data, the estimation pipeline and every number quoted in this paper are reproduced by a single script; see the reproducibility statement at the end.

Calibrated quantities. On the estimation window the implied risk aversion of the S&P 500 excess return series is $\delta = 2.44$, the ratio of mean to variance of monthly excess returns (the annualisation factor cancels); the S&P 500 serves as the proxy for δ while the equilibrium weights come from Idzorek’s asset classes, a mismatch we accept as the customary approximation. We use $\tau = 0.05$, the upper end of the range in He and Litterman (1999) and Idzorek (2007) (the central result of Section 6 is independent of this choice), and view confidences $c_k = 0.50$ throughout, an agnostic midpoint; the c -dependence of the leakage is explicit in the factor $c(1-c)$ of (3). The average asset variance, which acts as the Ledoit–Wolf shrinkage target scale, is $\mu_s = \frac{1}{n} \text{tr}(\hat{\Sigma}_s) = 0.0282$ annualised ($2.35 \cdot 10^{-3}$ per month), and the analytic shrinkage intensity is $\hat{\alpha} = 0.0406$, reported as $\alpha = 0.041$ throughout. Three relative views with these confidences run through the multi-view sections: IWF > IWD by +3% p.a., EEM > EFA by +4% p.a., and IWF > AGG by +6% p.a. The single-view analysis of Section 6 uses the first of these.

3 The Sample Covariance Problem

The sample covariance estimator is

$$\hat{\Sigma}_s = \frac{1}{T-1} \sum_{t=1}^T (r_t - \bar{r})(r_t - \bar{r})^\top.$$

For n assets and T observations its expected estimation loss grows with n/T (Ledoit and Wolf, 2004a). Once n/T is not negligible, the eigenvalues of $\hat{\Sigma}_s$ spread too widely; large eigenvalues are overestimated, small ones are underestimated (Ledoit and Wolf, 2004a, p. 371). The smallest eigenvalue moves toward zero, the condition number $\kappa(\Sigma) = \lambda_{\max}/\lambda_{\min}$ grows, and the inverse $\hat{\Sigma}_s^{-1}$ used in $w_{\text{BL}} = \frac{1}{\delta} \Sigma^{-1} \mu_{\text{BL}}$ amplifies estimation error. On the estimation window the condition number is $\kappa(\hat{\Sigma}_s) = 341.8$.

¹Some implementations map the posterior to weights with $\Sigma_{\text{post}} = \Sigma + M$, where M is the posterior covariance of the mean (He and Litterman, 1999). All results and all reported numbers in this paper use Σ itself; for relative views this keeps the weights exactly fully invested.

4 Ledoit–Wolf Shrinkage

Ledoit and Wolf (2004a) regularise $\hat{\Sigma}_s$ by shrinking it toward the scaled identity $\mu_s I_n$, with $\mu_s = \frac{1}{n} \text{tr}(\hat{\Sigma}_s)$,

$$\hat{\Sigma}_{\text{LW}} = (1 - \alpha) \hat{\Sigma}_s + \alpha \mu_s I_n,$$

where the oracle coefficient $\alpha \in [0, 1]$ minimises $\mathbb{E}[\|(1 - \alpha)\hat{\Sigma}_s + \alpha \mu_s I_n - \Sigma\|_F^2]$, with μ the population counterpart of μ_s ; the analytic estimator replaces the unobservable moments by consistent plug-ins, yielding $\hat{\alpha}$. Because $\alpha \mu_s I_n$ adds $\alpha \mu_s > 0$ to every eigenvalue of $(1 - \alpha)\hat{\Sigma}_s$, $\hat{\Sigma}_{\text{LW}} \succ 0$ is automatic. Throughout, we estimate $\hat{\alpha}$ with the analytic implementation of `scikit-learn` and apply it to the $T - 1$ -normalised sample matrix, so that the identity above holds exactly in the empirical pipeline. (The covariance output of `scikit-learn` itself normalises by T and would rescale every closed form below by $(T - 1)/T$.) On the estimation window the analytic estimator yields $\hat{\alpha} = 0.0406$ and brings the condition number down to $\kappa(\hat{\Sigma}_{\text{LW}}) = 108.0$, a factor 3.2 reduction.

5 Idzorek Ω : a Single View Is Inert

When Ω is derived from Σ via Idzorek’s calibration, the estimator cancels exactly in the view channel of a single view.

Idzorek’s confidence calibration leads, in the closed form derived by Walters (2014), to

$$\omega_{kk} = \frac{1 - c_k}{c_k} \tau p_k^\top \Sigma p_k$$

(Idzorek’s original procedure matches weight tilts iteratively; under the weight mapping of Section 2 the view tilt equals c_{eff} times the full-confidence tilt by (6), so matching a fraction c of it is solved exactly by this ω_{kk} in the single-view case). For a single view p the posterior then simplifies along the view direction to

$$p^\top \mu_{\text{BL}} = (1 - c) p^\top \pi + c Q,$$

an identity we derive in Section 6.6. The view absorption equals c exactly, independent of Σ . Switching $\hat{\Sigma}_s \rightarrow \hat{\Sigma}_{\text{LW}}$ scales the view variance and the matching ω proportionally, the ratio stays at c . Along the view direction the remaining LW impact runs through the prior spread $p^\top \pi$ alone; off the view direction the projection $\Sigma p / v_k$ moves as well. The prior shift is exactly linear in α ; the projection is a smooth rational function of α , linear to first order (and exactly linear only when p is an eigenvector of $\hat{\Sigma}_s$).

Concretely, the LW estimator shifts the prior linearly,

$$\Delta \pi_i = \alpha \delta \left[(\mu_s - \sigma_i^2) w_{\text{eq},i} - \sum_{j \neq i} \hat{\Sigma}_{s,ij} w_{\text{eq},j} \right].$$

Writing $\sigma_i^2 := \hat{\Sigma}_{s,ii}$, the diagonal term raises the prior of below-average-variance assets, but the cross-covariance term dominates in equity-heavy universes: the shift collapses to $\Delta \pi_i = \alpha \delta [\mu_s w_{\text{eq},i} - \text{Cov}(r_i, r_p)]$ with $r_p = w_{\text{eq}}^\top r$, so assets that co-move strongly with the equilibrium portfolio lose prior return. On our universe all six equity ETFs have $\Delta \pi_i < 0$ (including IWF and IWD, whose variances lie below μ_s), the two bond ETFs gain, and $|\Delta \pi_i| \leq 0.22\%$ p.a. throughout.

Remark 1 (several views). For $K > 1$ the absorption matrix is $A = S(S + \Omega)^{-1}$ with $S = \tau P \Sigma P^\top$. Idzorek’s Ω is tied to $\text{diag}(S)$ alone (up to the factor $(1 - c_k)/c_k$), so off-diagonal elements of S break both the exact absorption c and its estimator-invariance. On the three-view set of Section 2 the residuals $P \mu_{\text{BL}} - [(1 - c)P\pi + cQ]$ stay below 0.6% p.a., and the

per-view absorption remains within one percentage point of $c = 50\%$ and moves by less than 0.3 percentage points between the two estimators, an order of magnitude below the fixed- Ω shifts of Section 6.

This settles the Idzorek case for a single view; the fixed- Ω case does not cancel.

6 Fixed Ω : the Leakage Mechanism

A second workflow computes ω_k once on the sample estimator and stores the scalars. Later, the analyst replaces the covariance with an LW estimator, perhaps because the badly conditioned $\hat{\Sigma}_s^{-1}$ produced implausible leverage, but leaves ω_k fixed. We are not aware of survey evidence on how widespread this workflow is, though interfaces that accept a user-supplied Ω while Σ is estimated elsewhere make it easy to adopt; the point of this section is that following it changes the model's effective inputs without any visible change at the input stage. The view absorption shifts implicitly: the view variance v_k (defined below) moves with the estimator while ω_k stays fixed.

We derive the closed form for the shift in c_{eff} , translate it into a general sign-and-size criterion and, for the symmetric case, into a critical correlation threshold, and verify the chain on the IWF/IWD pair. Throughout this section the analysis rests on a single view, taken as the relative pair $p = e_i - e_j$ unless stated otherwise; the symmetric per-component statements additionally assume $\sigma_i \approx \sigma_j$ for the view pair, an assumption we make explicit where it is used and drop where it is not needed.

6.1 Effective View Absorption c_{eff}

The actual view absorption after evaluating the posterior is

$$c_{\text{eff},k} = \frac{\tau v_k}{\tau v_k + \omega_k}, \quad v_k := p_k^\top \Sigma p_k.$$

$c_{\text{eff},k}$ measures how far the posterior travels along the view direction from prior toward view. At $c_{\text{eff},k} = 0$ the posterior stays at the prior, at $c_{\text{eff},k} = 1$ it lands exactly on the view. The ratio depends on the view variance v_k , the view uncertainty ω_k , and the prior scale τ , which controls how tightly the prior is held: at fixed ω_k , small τ pulls $c_{\text{eff},k}$ down, large τ lets the posterior move toward the view more easily.

In the Idzorek world with $\omega_k = \frac{1-c_k}{c_k} \tau v_k$, both τ and v_k cancel and $c_{\text{eff},k} = c_k$ exactly. In the fixed- Ω world this cancellation breaks.

6.2 Shift of v_k Under Shrinkage

For a relative view $p_k = e_i - e_j$ the view variance is $v_k = \sigma_i^2 + \sigma_j^2 - 2\hat{\Sigma}_{s,ij} = \sigma_i^2 + \sigma_j^2 - 2\rho_{ij}\sigma_i\sigma_j$, with no assumption on the volatilities. Under LW,

$$v_k^{\text{LW}} = (1 - \alpha) v_k^{\text{Std}} + 2\alpha \mu_s,$$

since $p_k^\top I_n p_k = 2$. Hence v_k grows under shrinkage exactly when $v_k^{\text{Std}} < 2\mu_s$, i.e. when the view variance lies below twice the average asset variance. For positively correlated pairs of similar volatility this is the typical case; it also holds for pairs whose variances are simply small relative to the universe average, such as bond pairs. At fixed ω_k , the absorption $c_{\text{eff},k}$ grows with v_k .

The pair structure of the view enters only through $p_k^\top I_n p_k$. For an arbitrary view vector p_k ,

$$v_k^{\text{LW}} = (1 - \alpha) v_k^{\text{Std}} + \alpha \mu_s \|p_k\|^2,$$

and every formula below holds verbatim with $2\mu_s$ replaced by $\mu_s \|p_k\|^2$: the general sign criterion is $v_k^{\text{Std}} \leq \mu_s \|p_k\|^2$. Two special cases follow immediately. An *absolute* view on asset i ($p_k = e_i$,

$\|p_k\|^2 = 1$) inflates exactly when $\sigma_i^2 < \mu_s$: on our universe a frozen- ω absolute view on AGG shifts by +12.9 percentage points, the largest leakage in the universe, ahead of every pair, while an absolute view on EEM ($\sigma_{\text{EEM}}^2 > \mu_s$) shifts by -0.5 percentage points, the deflation regime realised in the data. A basket view (equal-weighted US growth over US value, $\|p_k\|^2 = 1$) shifts by +5.8 percentage points.

6.3 Closed Form for $\Delta c_{\text{eff},k}$

Combining $c_{\text{eff},k}(\hat{\Sigma}_{\text{LW}})$ and $c_{\text{eff},k}(\hat{\Sigma}_s)$ over a common denominator and simplifying,

$$\begin{aligned}\Delta c_{\text{eff},k} &= \frac{\tau v_k^{\text{LW}}}{\tau v_k^{\text{LW}} + \omega_k} - \frac{\tau v_k^{\text{Std}}}{\tau v_k^{\text{Std}} + \omega_k} \\ &= \frac{\tau v_k^{\text{LW}} (\tau v_k^{\text{Std}} + \omega_k) - \tau v_k^{\text{Std}} (\tau v_k^{\text{LW}} + \omega_k)}{(\tau v_k^{\text{LW}} + \omega_k)(\tau v_k^{\text{Std}} + \omega_k)} \\ &= \frac{\tau \omega_k (v_k^{\text{LW}} - v_k^{\text{Std}})}{(\tau v_k^{\text{LW}} + \omega_k)(\tau v_k^{\text{Std}} + \omega_k)}.\end{aligned}$$

Substituting $v_k^{\text{LW}} - v_k^{\text{Std}} = \alpha(2\mu_s - v_k^{\text{Std}})$ yields

$$\Delta c_{\text{eff},k} = \frac{\tau \omega_k \alpha (2\mu_s - v_k^{\text{Std}})}{(\tau v_k^{\text{LW}} + \omega_k)(\tau v_k^{\text{Std}} + \omega_k)}. \quad (2)$$

For small α the shift is linear in α , and its sign matches $\text{sgn}(2\mu_s - v_k^{\text{Std}})$ for any α : the general criterion announced above, valid for arbitrary volatility ratios.

In the workflow under study, ω_k was calibrated on the sample estimator, $\omega_k = \frac{1-c}{c} \tau v_k^{\text{Std}}$. Substituting this into (2) cancels τ completely and gives the exact form

$$\Delta c_{\text{eff},k} = \frac{c(1-c)(v_k^{\text{LW}} - v_k^{\text{Std}})}{c v_k^{\text{LW}} + (1-c)v_k^{\text{Std}}}. \quad (3)$$

Three consequences are worth recording. First, the leakage does not depend on the prior scale τ at all; no choice of τ mitigates it. Second, (3) is exact for any pair of positive view variances, so it can be evaluated asset pair by asset pair without any symmetry assumption. Third, at fixed variances the shift is maximised over the confidence at $c^* = 1/(1 + \sqrt{v_k^{\text{LW}}/v_k^{\text{Std}}})$; on the IWF/IWD pair $c^* = 0.45$ with a shift of +10.9 percentage points, so the agnostic $c = 0.50$ used throughout operates close to the worst case. As $v_k^{\text{Std}} \rightarrow 0$ at fixed μ_s the shift approaches its supremum $1 - c$.

Remark 2 (range and boundary). As a function of $v_k^{\text{Std}} \in (0, \infty)$ at fixed target $\mu_s \|p_k\|^2$, the shift (3) is strictly decreasing (the numerator of its derivative is $-c(1-c)\alpha\mu_s \|p_k\|^2$), with range

$$\Delta c_{\text{eff},k} \in \left(-\frac{c(1-c)\alpha}{1-c\alpha}, 1-c \right).$$

The leakage is therefore asymmetric: the inflation end is bounded only by $1 - c$ (fifty percentage points at $c = 0.50$), while the deflation end has an $O(\alpha)$ floor: at $\alpha = 0.041$ and $c = 0.50$ no view in any universe can lose more than 1.05 percentage points of effective confidence, and even at $\rho_{ij} = -1$ in the symmetric setting of Section 6.4 the shift is -0.4 ($\alpha = 0.041$) to -3.4 percentage points ($\alpha = 0.30$). Both endpoints are boundary limits: $v_k^{\text{Std}} \rightarrow 0$ corresponds to $\rho_{ij} \rightarrow 1$ for the equal-volatility pair view, and $\rho_{ij} = \pm 1$ itself lies on the boundary of the positive-definite cone, where the posterior machinery degenerates but (3), which needs only $v_k^{\text{Std}} > 0$, extends continuously. Note also that shrinkage enforces $v_k^{\text{LW}} \geq \alpha\mu_s \|p_k\|^2 > 0$: Ledoit-Wolf regularises exactly the near-degenerate view directions in which the frozen- ω confidence explodes.

6.4 Critical Correlation ρ^* in the Symmetric Case

To turn (3) into a rule of thumb, consider the symmetric case $\sigma_i = \sigma_j = \sigma$, in which $v_k^{\text{Std}} = 2\sigma^2(1 - \rho_{ij})$. Setting $\Delta c_{\text{eff},k} = \varepsilon$ in (3) and solving for ρ_{ij} with $v_k^{\text{LW}} = (1 - \alpha)v_k^{\text{Std}} + 2\alpha\mu_s$ yields the closed-form threshold

$$\rho^*(\alpha, \varepsilon) = 1 - \frac{c\alpha\mu_s[(1 - c) - \varepsilon]}{\sigma^2[\varepsilon(1 - c\alpha) + c(1 - c)\alpha]}, \quad \varepsilon \in (0, 1 - c). \quad (4)$$

A view triggers $\Delta c_{\text{eff},k} > \varepsilon$ exactly when $\rho_{ij} > \rho^*$. The threshold falls as α rises and as the tolerated shift ε falls. (On the negative branch, $\Delta c_{\text{eff},k} < -\varepsilon$ requires ρ_{ij} below a second threshold ρ^{**} , which exists once $\varepsilon < c(1 - c)\alpha/(1 - c\alpha)$; at $\alpha = 0.30$ and $\varepsilon = 1\%$ it sits at $\rho^{**} = -0.32$, at $\alpha = 0.041$ it lies far outside $[-1, 1]$. The tables below concern the positive branch; Remark 2 bounds the entire negative branch.) The symmetric case serves as a rule of thumb: for $\sigma_i \neq \sigma_j$ the exact criterion remains (3) evaluated at the heterogeneous $v_k^{\text{Std}} = \sigma_i^2 + \sigma_j^2 - 2\rho_{ij}\sigma_i\sigma_j$, and we use that exact form for the pair-level results below.

6.5 Values and Economic Statement

Table 1 shows ρ^* for four shrinkage intensities and four tolerance levels ε . The parameters $\mu_s = 4.0 \cdot 10^{-3}$, $\sigma^2 = 3.5 \cdot 10^{-3}$, $c = 0.50$ are illustrative, chosen to match the order of magnitude of the monthly quantities in our data ($\text{tr}(\hat{\Sigma}_s)/n = 2.35 \cdot 10^{-3}$ per month); the table depends on the parameters only through the ratio μ_s/σ^2 and c .

Table 1: Critical correlation $\rho^*(\alpha, \varepsilon)$ from (4). The pair correlation ρ_{ij} must exceed this value for $\Delta c_{\text{eff}} > \varepsilon$. Illustrative parameters $\mu_s = 4.0 \cdot 10^{-3}$, $\sigma^2 = 3.5 \cdot 10^{-3}$, $c = 0.50$.

ε	$\alpha = 0.04$	$\alpha = 0.10$	$\alpha = 0.166$	$\alpha = 0.30$
1 %	0.434	0.188	0.083	-0.006
2 %	0.629	0.377	0.239	0.106
5 %	0.826	0.645	0.511	0.343
10 %	0.915	0.810	0.715	0.571

At moderate shrinkage the threshold is demanding only for tight tolerances: with $\alpha = 0.04$, a five-percentage-point shift already occurs above $\rho^* = 0.826$, a correlation level typical of style pairs within one equity market ($\rho = 0.920$ for IWF/IWD, Table 2). At strong shrinkage ($\alpha = 0.30$), every non-negatively correlated pair crosses the one-percent band. The analytic intensity tends to grow with n/T (Ledoit and Wolf, 2004a), so larger universes on the same window typically sit further to the right in the table.

The symmetric threshold understates the reach of the effect, because the general criterion is $v_k^{\text{Std}} < 2\mu_s$, not high correlation per se. Table 2 evaluates the exact expression (3) for every pair of the universe, one view at a time, with the full 8×8 covariance shrunk once at $\alpha = 0.041$ (the practitioner workflow) and $c = 0.50$; the table reports the pairs above the five-point band. Five of 28 pairs shift by more than five percentage points. The top of the list is dominated by the highly correlated, nearly homogeneous style pairs, led by IWF/IWD at +10.7 percentage points: the model operates at an effective confidence of 0.61 where the analyst specified 0.50. But the bond pair AGG/BWX ranks third at +8.2 percentage points despite a moderate correlation of 0.57: its view variance is small relative to $2\mu_s$ because bond variances are far below the universe average. A threshold stated in correlations alone would miss it. In this universe every one of the 28 pairs satisfies $v_k^{\text{Std}} < 2\mu_s$, so the leakage is a one-sided confidence inflation; the deflation regime $v_k^{\text{Std}} > 2\mu_s$ requires pairs whose spread variance exceeds twice the universe average, such as weakly correlated high-volatility assets.

Table 2: Exact Δc_{eff} from (3) for all pairs with a shift above five percentage points; each pair is evaluated as a single relative view in isolation. Full-universe shrinkage at $\alpha = 0.041$, ω frozen on the sample estimator, $c = 0.50$; $k = \sigma_i/\sigma_j$. Estimation window 2008–2018.

Pair	ρ_{ij}	k	Δc_{eff} (pp)
IWF/IWD	0.920	0.99	+10.7
IWO/IWN	0.939	1.03	+8.8
AGG/BWX	0.574	0.45	+8.2
IWD/IWN	0.917	0.80	+6.6
IWF/IWO	0.909	0.77	+5.5

In economic terms, whenever a relative view connects two assets whose spread variance is small (because they are highly correlated, or because both carry little variance to begin with), a frozen Ω lets the shrinkage-induced widening of the spread variance flow into extra view confidence, at the rate given by (3). The effective confidence exceeds the stated confidence. For weakly correlated, high-variance pairs or negligible α , c_{eff} stays close to c and the estimator switch is harmless.

6.6 From Δc_{eff} to $\Delta\mu$ and Δw

The shifts $\Delta\mu$ and Δw follow analytically from Δc_{eff} . The key step is an exact closed form for the posterior under a single relative view, which we derive directly from the posterior equation.

With a single view, $P = p^\top$ and $\Omega = \omega$ is a scalar; the BL posterior equation reads

$$M \mu_{\text{BL}} = (\tau\Sigma)^{-1}\pi + \frac{1}{\omega} p Q, \quad M = (\tau\Sigma)^{-1} + \frac{1}{\omega} p p^\top.$$

Left-multiplication by $\tau\Sigma$ and solving for μ_{BL} :

$$\begin{aligned} \tau\Sigma M \mu_{\text{BL}} &= \tau\Sigma [(\tau\Sigma)^{-1}\pi + \frac{1}{\omega} p Q] \\ \left[I + \frac{\tau}{\omega} \Sigma p p^\top\right] \mu_{\text{BL}} &= \pi + \frac{\tau Q}{\omega} \Sigma p \\ \mu_{\text{BL}} + \frac{\tau(p^\top \mu_{\text{BL}})}{\omega} \Sigma p &= \pi + \frac{\tau Q}{\omega} \Sigma p \\ \mu_{\text{BL}} &= \pi + \frac{\tau(Q - p^\top \mu_{\text{BL}})}{\omega} \Sigma p. \end{aligned} \tag{*}$$

Applying $p^\top$ to (*) produces a scalar equation in $p^\top \mu_{\text{BL}}$ with $v_k := p^\top \Sigma p$:

$$\begin{aligned} p^\top \mu_{\text{BL}} &= p^\top \pi + \frac{\tau v_k (Q - p^\top \mu_{\text{BL}})}{\omega} \\ \omega p^\top \mu_{\text{BL}} &= \omega p^\top \pi + \tau v_k Q - \tau v_k p^\top \mu_{\text{BL}} \\ (\omega + \tau v_k) p^\top \mu_{\text{BL}} &= \omega p^\top \pi + \tau v_k Q \\ p^\top \mu_{\text{BL}} &= \frac{\omega}{\omega + \tau v_k} p^\top \pi + \frac{\tau v_k}{\omega + \tau v_k} Q \\ &= (1 - c_{\text{eff}}) p^\top \pi + c_{\text{eff}} Q, \quad c_{\text{eff}} := \frac{\tau v_k}{\tau v_k + \omega}. \end{aligned}$$

Hence $Q - p^\top \mu_{\text{BL}} = (1 - c_{\text{eff}})(Q - p^\top \pi)$. Inserting this into (*) and simplifying the prefactor:

$$\begin{aligned} \mu_{\text{BL}} &= \pi + \frac{\tau(1 - c_{\text{eff}})}{\omega} (Q - p^\top \pi) \Sigma p \\ &= \pi + \frac{\tau}{\omega + \tau v_k} (Q - p^\top \pi) \Sigma p \\ &= \pi + c_{\text{eff}} (Q - p^\top \pi) \frac{\Sigma p}{v_k}. \end{aligned}$$

The exact identity is established:

$$\mu_{\text{BL}} = \pi + c_{\text{eff}} (Q - p^\top \pi) \frac{\Sigma p}{v_k}, \quad c_{\text{eff}} = \frac{\tau v_k}{\tau v_k + \omega}. \tag{5}$$

The posterior moves exactly along $\Sigma p/v_k$, normalised so that $p^\top(\Sigma p/v_k) = 1$. Substituting (5) into $w_{\text{BL}} = \frac{1}{\delta}\Sigma^{-1}\mu_{\text{BL}}$ simplifies the weight expression as well,

$$w_{\text{BL}} = w_{\text{eq}} + c_{\text{eff}}(Q - p^\top\pi)\frac{p}{\delta v_k}, \quad w_{\text{eq}} = \frac{1}{\delta}\Sigma^{-1}\pi. \quad (6)$$

The view-driven part of w_{BL} lies exactly in direction p , without any assumption on the volatilities. Equation (6) is the single-view case of the unconstrained-weights result of [He and Litterman \(1999\)](#) (posterior weights are the market weights plus a position in the view portfolio) and serves here as the carrier that propagates Δc_{eff} .

Step 1: $\Delta\mu$ from Δc_{eff} . Projecting (5) on p gives

$$p^\top \Delta\mu_{\text{BL}} = \Delta c_{\text{eff}}(Q - p^\top\pi),$$

exact for any pair of estimators, since $p^\top(\Sigma p/v_k) = 1$ under each. The full vector shift additionally moves the direction $\Sigma p/v_k$; holding Σ fixed, it reduces to

$$\Delta\mu_{\text{BL}} = \Delta c_{\text{eff}}(Q - p^\top\pi)\frac{\Sigma p}{v_k}.$$

For $p = e_i - e_j$ with $\sigma_i \approx \sigma_j$ one has $\Sigma p \approx (v_k/p^\top p)p$, hence

$$\Delta\mu_{\text{BL}} \approx \frac{\Delta c_{\text{eff}}(Q - p^\top\pi)}{p^\top p}p, \quad \Delta\mu_i \approx -\Delta\mu_j \approx \frac{1}{2}\Delta c_{\text{eff}}(Q - p^\top\pi). \quad (7)$$

This symmetric per-component split rests entirely on $\sigma_i \approx \sigma_j$. If the volatilities differ, the direction $\Sigma p/v_k$ distributes the return shift asymmetrically across the two assets and the clean antisymmetric split no longer holds; the exact vector identity (5) is unaffected, only its symmetric reading depends on the assumption. The numerical example below uses a pair with nearly identical volatilities and confirms that the symmetric split is accurate in the homogeneous case.

Step 2: Δw from $\Delta\mu$. Differentiating (6) with respect to c_{eff} at fixed Σ yields directly

$$\Delta w_{\text{BL}} = \frac{\Delta c_{\text{eff}}(Q - p^\top\pi)}{\delta v_k}p. \quad (8)$$

The shift sits exactly along p with per-component $\Delta w_i = -\Delta w_j = \Delta c_{\text{eff}}(Q - p^\top\pi)/(\delta v_k)$. (8) is the exact partial derivative at fixed Σ , i.e. the pure view-absorption channel. The 2×2 example below, which uses $\hat{\Sigma}_s^{-1}$ for both scenarios, measures it up to the small direction shift $\Sigma_{\text{LW}}p/v_k^{\text{LW}} - \hat{\Sigma}_s p/v_k^{\text{Std}}$ that the shared inverse does not remove. The separate effect of a changed Σ^{-1} is the precision channel from Section 7.

Every percentage point of Δc_{eff} moves the spread $w_i - w_j$ linearly, with leverage $1/(\delta v_k)$. Small v_k amplifies the effect, and small v_k is exactly the regime in which Δc_{eff} is large. Affected pairs therefore combine the largest confidence drift with the strongest translation factor into the weights.

6.7 A Worked 2×2 Example: IWF and IWD

We isolate the effect on the expected-return vector by reducing the universe to two assets, IWF (US large-cap growth) and IWD (US large-cap value), and running the BL formula in closed form. No other assets interact, so every number below is fully traceable. All quantities are annualised and refer to the estimation window January 2008 to December 2018. With $\sigma_{\text{IWF}} = 15.5\%$ and $\sigma_{\text{IWD}} = 15.6\%$ ($k = \sigma_i/\sigma_j = 0.99$) the pair lies in the homogeneous special case, so the symmetric approximations of Section 6.6 apply.

To keep the example self-contained, we shrink the extracted 2×2 sample block with its own target $\mu_s^{(2)} = \frac{1}{2}\text{tr}(\hat{\Sigma}_s^{(2)})$, while the intensity $\alpha = 0.041$ is the analytic value from the full eight-asset universe. (If we shrink the full 8×8 matrix instead and extract the pair afterwards, the design underlying Table 2, the larger universe-wide target $\mu_s = 0.0282 > \mu_s^{(2)} = 0.0242$ widens the spread variance further and the shift grows to +10.7 percentage points; the self-contained design is thus the conservative one.)

The inputs are

$$\delta = 2.44, \quad \tau = 0.05, \quad c = 0.50, \quad w_{\text{eq}} = (0.50, 0.50)^\top, \quad Q = 0.03.$$

The two annualised covariance matrices read

$$\begin{aligned} \hat{\Sigma}_s &= \begin{pmatrix} 0.02408 & 0.02227 \\ 0.02227 & 0.02435 \end{pmatrix}, & \rho_s &= 0.920, \\ \hat{\Sigma}_{\text{LW}} &= \begin{pmatrix} 0.02409 & 0.02137 \\ 0.02137 & 0.02435 \end{pmatrix}, & \rho_{\text{LW}} &= 0.882, & \alpha &= 0.041. \end{aligned}$$

We calibrate ω once on the sample estimator and freeze it, $\omega = \frac{1-c}{c} \tau v_k^{\text{Std}} = 1.95 \cdot 10^{-4}$.

The effective view absorption follows directly,

$$\begin{aligned} v_k^{\text{Std}} &= 0.00389, & v_k^{\text{LW}} &= 0.00570, \\ c_{\text{eff}}(\hat{\Sigma}_s) &= 50.00 \%, & c_{\text{eff}}(\hat{\Sigma}_{\text{LW}}) &= 59.42 \%, \\ & & \Delta c_{\text{eff}} &= +9.42 \text{ pp}. \end{aligned}$$

The expected-return vector and the corresponding portfolio weights shift accordingly. The LW weights deliberately use $\hat{\Sigma}_s^{-1}$ as well, so only the view-absorption channel shows up; the separate precision channel through $\hat{\Sigma}^{-1}$ is treated in Section 7.

$$\begin{aligned} \mu_{\text{BL}}^{\text{Std}} &= (6.37 \%, 4.89 \%), & \mu_{\text{BL}}^{\text{LW}} &= (6.52 \%, 4.75 \%), \\ w_{\text{BL}}^{\text{Std}} &= (+209.4 \%, -109.4 \%), & w_{\text{BL}}^{\text{LW}} &= (+239.6 \%, -139.2 \%). \end{aligned}$$

The difference vectors read

$$\boxed{\Delta c_{\text{eff}} = +9.42 \text{ pp}, \quad \Delta \mu = (+0.15 \text{ pp}, -0.13 \text{ pp}), \quad \Delta w = (+30.2 \text{ pp}, -29.8 \text{ pp}).}$$

The closed-form formulas from the previous subsection reproduce these numbers. With $p^\top p = 2$, $Q - p^\top \pi = 0.0303$, $\delta = 2.44$ and $v_k = 0.00389$ from the sample estimator we check

$$\begin{aligned} \Delta \mu_{\text{IWF}} &\approx \frac{1}{2} \Delta c_{\text{eff}} (Q - p^\top \pi) = \frac{1}{2} \cdot 0.0942 \cdot 0.0303 = +0.143 \text{ pp}, \\ \Delta w_{\text{IWF}} &\approx \frac{\Delta c_{\text{eff}} (Q - p^\top \pi)}{\delta v_k} = \frac{0.0942 \cdot 0.0303}{2.44 \cdot 0.00389} \approx +0.300 = +30.0 \text{ pp}. \end{aligned}$$

The IWD counterparts have the opposite sign. The $\Delta \mu$ prediction matches the numerical value to within 0.01 pp, the Δw prediction to within 0.2 pp; both gaps stem from the slight asymmetry $\sigma_{\text{IWF}} \neq \sigma_{\text{IWD}}$ that the approximation $\Sigma p / v_k \approx p / (p^\top p)$ ignores.

The chain runs from the correlation drop to the weights: the correlation between IWF and IWD falls from 0.920 to 0.882 under shrinkage, v_k grows from 0.00389 to 0.00570, and c_{eff} climbs by 9.42 percentage points. The posterior pushes IWF up by 0.15 pp and IWD down by 0.13 pp. Through $w = \frac{1}{\delta} \Sigma^{-1} \mu$ this return shift translates into roughly ± 30 percentage points on the long-short spread. The analyst's input form still shows $c = 50 \%$; the model operates at $c_{\text{eff}} = 59.42 \%$. The absolute position sizes illustrate a known property of unconstrained Black-Litterman with highly correlated pairs: the translation factor $1/(\delta v_k) \approx 105$ turns any view into triple-digit tilts, and the leakage is amplified by the same factor.

6.8 Recommendation

The clean fix is to recompute ω_k alongside Σ ,

$$\omega_k^{\text{LW}} = \frac{1 - c_k}{c_k} \tau v_k^{\text{LW}}.$$

This restores $c_{\text{eff},k} = c_k$ exactly. The same self-healing holds for any Ω tied functionally to Σ , for instance the He–Litterman convention $\Omega = \text{diag}(P \tau \Sigma P^\top)$: the leakage is a property of frozen scalars, independent of the particular calibration formula. Covariance estimator, confidence and view-uncertainty form one unit. Replacing one of the three pieces without re-deriving the other two is the source of the leakage, not the choice of estimator itself. An analyst who freezes Ω has raised the confidence of every affected view without a decision to do so, in conflict with the risk-governance process that produced the original confidences.

One caveat delimits the recommendation. It presumes that the analyst’s primitive is the stated confidence c_k , as it is under Idzorek’s procedure and under any sign-off process phrased in confidences. If the primitive is instead ω_k itself, an absolute variance attached to the view noise irrespective of the covariance model, then freezing ω_k is Bayesianly coherent (cf. [Kolm and Ritter, 2017](#)), and the induced change in $c_{\text{eff},k}$ is the correct posterior response to a better estimate of v_k ; recomputing Ω would then be the error. The two readings prescribe opposite fixes, which is why the choice of primitive belongs in the model documentation. Throughout this paper we take c_k as the primitive.

7 Conditioning and Precision Channel

The Idzorek case showed that the view absorption stays stable. The fixed- Ω case showed that the shift in the return signal becomes material under identifiable conditions. The larger effect of the LW switch lies in the inverse $\hat{\Sigma}^{-1}$ entering $w_{\text{BL}} = \frac{1}{\delta} \Sigma^{-1} \mu_{\text{BL}}$. The weight change decomposes into two channels,

$$\Delta w_{\text{BL}} = \underbrace{\frac{1}{\delta} \hat{\Sigma}_{\text{LW}}^{-1} (\mu_{\text{BL}}^{\text{LW}} - \mu_{\text{BL}}^{\text{Std}})}_{\text{return channel}} + \underbrace{\frac{1}{\delta} (\hat{\Sigma}_{\text{LW}}^{-1} - \hat{\Sigma}_s^{-1}) \mu_{\text{BL}}^{\text{Std}}}_{\text{precision channel}}.$$

The return channel is small because Idzorek fixes the view absorption at c ; the remaining terms, the prior shift and the projection shift, are both $O(\alpha)$. Under the three-view set of Section 2 the largest change in any posterior mean across the two estimators is 0.42 % p.a., of which the prior channel contributes at most 0.21 % (Section 5); the remainder enters through the Σ -dependent view projection at nearly unchanged absorption (Remark 1). The precision channel is driven by the difference of inverses.

LW lifts every eigenvalue λ_k of $\hat{\Sigma}_s$,

$$\lambda_k^{\text{LW}} = (1 - \alpha) \lambda_k + \alpha \mu_s, \quad k = 1, \dots, n.$$

The smallest eigenvalue rises by the largest factor, so $1/\lambda_n$ falls the most; on our data it falls by roughly a factor of three. As a result $\hat{\Sigma}_{\text{LW}}^{-1}$ is much more uniform than $\hat{\Sigma}_s^{-1}$, and the LW inverse pulls back the extreme long-short positions that the badly conditioned estimator produces along low-variance directions. On our universe the condition number drops from 341.8 to 108.0; under the three-view set of Section 2 (weights via $w = \frac{1}{\delta} \Sigma^{-1} \mu_{\text{BL}}$), the IWF position (US large-cap growth) falls from +173.5 % to +121.0 % and the IWD short (US large-cap value) from −136.4 % to −83.0 %, taking the gross exposure from 4.31 to 3.16.

The cost of the frozen- Ω workflow appears at the same level, and it requires no data beyond the estimation window. With Ω frozen on the sample estimator, the inflated view confidences partially restore the trimmed positions (IWF +140.5 % instead of +121.0 %; gross exposure 3.63

instead of 3.16) and undo roughly 40 % of the exposure reduction that the recomputed- Ω switch delivers. Economically, the inflated confidence acts as leverage on the view: the position scales with c_{eff} , and with it both the payoff if the view is right and the loss if it is wrong. A performance evaluation on subsequent data would add only a single realised path of this trade-off; we therefore report the weight-level cost.

8 Conclusion

Within the view channel, whether the Ledoit–Wolf switch is harmless inside Black–Litterman is decided by the treatment of Ω , and both regimes now have closed forms. Under Idzorek’s calibration $\omega_k(\Sigma) = \frac{1-c_k}{c_k} \tau p_k^\top \Sigma p_k$ the absorption $c_{\text{eff},k}$ is independent of Σ for a single view: any positive-definite estimator can be plugged in without distorting the view channel, and the posterior impact is confined to prior and projection shifts, the former exactly linear and the latter first-order linear in α . If ω_k is instead computed once and frozen, the estimator switch shifts the effective confidence by the exact, τ -free expression (3). The sign criterion is $v_k^{\text{Std}} \leq \mu_s \|p_k\|^2$ ($2\mu_s$ for relative pairs), the shift is confined to the band $(-c(1-c)\alpha/(1-c\alpha), 1-c)$ of Remark 2, the symmetric special case yields the closed-form correlation threshold (4), and the exact single-view identity (5) propagates the shift into posterior means and weights with leverage $1/(\delta v_k)$. On the eight-asset universe at $\alpha = 0.041$, five of 28 pairs cross the five-percentage-point band, led by IWF/IWD at +10.7; the bond pair AGG/BWX crosses it at a correlation of only 0.57. In the weights, the frozen- Ω workflow restores roughly 40 % of the gross-exposure reduction that the estimator switch delivers. Recomputing ω_k alongside Σ removes the leakage analytically. The lesson is procedural: covariance, confidence and view uncertainty have to be kept consistent with each other.

Several limitations point to extensions. The closed-form chain covers a single view (arbitrary view vectors in the shift formula, relative pairs for the per-component weight split) and a diagonal Ω ; with several simultaneous views the shifts interact through the posterior precision (Remark 1 bounds the size of these cross-terms on our data), and the multi-view analogue (via Sherman–Morrison updates) we leave to future work. The empirical part rests on a single estimation window and a small, well-diversified ETF universe on which the analytic intensity is modest; equity universes with larger n/T push α , and with it the leakage, substantially higher. Finally, the mechanism is target-independent (replacing the scaled identity by the single-factor target of Ledoit and Wolf (2003) or the nonlinear shrinkage of Ledoit and Wolf (2017) only changes the expression for $v_k^{\text{LW}} - v_k^{\text{Std}}$), but the quantitative reach of the effect under these estimators remains to be mapped out.

Reproducibility

All numbers quoted in this paper are produced by the script `code/paper_numbers.py` from frozen monthly return data (`data/`, sourced from Yahoo Finance). Pipeline: excess returns; sample covariance with $T-1$ normalisation; the analytic Ledoit–Wolf intensity of `scikit-learn` applied to that matrix (Section 4); weights via $w = (\delta \Sigma)^{-1} \mu_{\text{BL}}$; implied risk aversion from S&P 500 excess returns. Code and data are available at <https://github.com/lucasposern/blm-lw-fixed-omega>.

References

W. Bessler, H. Opfer, and D. Wolff. Multi-asset portfolio optimization and out-of-sample performance: An evaluation of Black–Litterman, mean-variance, and naïve diversification approaches. *The European Journal of Finance*, 23(1):1–30, 2017.

- M. J. Best and R. R. Grauer. On the sensitivity of mean-variance-efficient portfolios to changes in asset means: Some analytical and computational results. *The Review of Financial Studies*, 4(2):315–342, 1991.
- A. Bevan and K. Winkelmann. Using the Black–Litterman global asset allocation model: Three years of practical experience. Fixed Income Research, Goldman Sachs & Co., 1998.
- F. Black and R. Litterman. Global portfolio optimization. *Financial Analysts Journal*, 48(5):28–43, 1992.
- V. K. Chopra and W. T. Ziemba. The effect of errors in means, variances, and covariances on optimal portfolio choice. *The Journal of Portfolio Management*, 19(2):6–11, 1993.
- V. DeMiguel, L. Garlappi, and R. Uppal. Optimal versus naive diversification: How inefficient is the $1/N$ portfolio strategy? *The Review of Financial Studies*, 22(5):1915–1953, 2009.
- V. DeMiguel, A. Martín-Utrera, and F. J. Nogales. Size matters: Optimal calibration of shrinkage estimators for portfolio selection. *Journal of Banking & Finance*, 37(8):3018–3034, 2013.
- P. A. Frost and J. E. Savarino. An empirical Bayes approach to efficient portfolio selection. *Journal of Financial and Quantitative Analysis*, 21(3):293–305, 1986.
- G. He and R. Litterman. The intuition behind Black–Litterman model portfolios. Investment Management Research, Goldman Sachs & Co., 1999.
- U. Herold. Portfolio construction with qualitative forecasts. *The Journal of Portfolio Management*, 30(1):61–72, 2003.
- T. Idzorek. A step-by-step guide to the Black–Litterman model: Incorporating user-specified confidence levels. In S. Satchell, editor, *Forecasting Expected Returns in the Financial Markets*, pages 17–38. Academic Press, 2007.
- R. Jagannathan and T. Ma. Risk reduction in large portfolios: Why imposing the wrong constraints helps. *The Journal of Finance*, 58(4):1651–1683, 2003.
- P. Jorion. Bayes–Stein estimation for portfolio analysis. *Journal of Financial and Quantitative Analysis*, 21(3):279–292, 1986.
- R. Kan and G. Zhou. Optimal portfolio choice with parameter uncertainty. *Journal of Financial and Quantitative Analysis*, 42(3):621–656, 2007.
- P. N. Kolm and G. Ritter. On the Bayesian interpretation of Black–Litterman. *European Journal of Operational Research*, 258(2):564–572, 2017.
- O. Ledoit and M. Wolf. Improved estimation of the covariance matrix of stock returns with an application to portfolio selection. *Journal of Empirical Finance*, 10(5):603–621, 2003.
- O. Ledoit and M. Wolf. A well-conditioned estimator for large-dimensional covariance matrices. *Journal of Multivariate Analysis*, 88(2):365–411, 2004a.
- O. Ledoit and M. Wolf. Honey, I shrunk the sample covariance matrix. *The Journal of Portfolio Management*, 30(4):110–119, 2004b.
- O. Ledoit and M. Wolf. Nonlinear shrinkage of the covariance matrix for portfolio selection: Markowitz meets Goldilocks. *The Review of Financial Studies*, 30(12):4349–4388, 2017.
- O. Ledoit and M. Wolf. Analytical nonlinear shrinkage of large-dimensional covariance matrices. *The Annals of Statistics*, 48(5):3043–3065, 2020.

- H. Markowitz. Portfolio selection. *The Journal of Finance*, 7(1):77–91, 1952.
- A. Meucci. The Black–Litterman approach: Original model and extensions. In R. Cont, editor, *Encyclopedia of Quantitative Finance*. Wiley, Chichester, 2010.
- R. O. Michaud. The Markowitz optimization enigma: Is ‘optimized’ optimal? *Financial Analysts Journal*, 45(1):31–42, 1989.
- R. O. Michaud, D. N. Esch, and R. Michaud. Deconstructing Black–Litterman: How to get the portfolio you already knew you wanted. *Journal of Investment Management*, 11(1):6–20, 2013.
- S. Satchell and A. Scowcroft. A demystification of the Black–Litterman model: Managing quantitative and traditional portfolio construction. *Journal of Asset Management*, 1(2):138–150, 2000.
- W. F. Sharpe. Capital asset prices: A theory of market equilibrium under conditions of risk. *The Journal of Finance*, 19(3):425–442, 1964.
- J. Walters. The Black–Litterman model in detail. Working paper, SSRN 1314585, 2014.